




P50 time to first token




P99 time to first token




Per-request mean TPOT percentiles



Output throughput (whole TP8 replica)



Cache hits under two timing policies



Observed concurrency peaks

